Paper 27 Integration into `scpn-control`

Kuramoto–Sakaguchi Phase Reduction with
Exogenous Global Field Driver $\zeta \sin(\Psi - \theta)$

Miroslav Šotek

ORCID: 0009-0009-3560-0851

www.anulum.li | protoscience@anulum.li

25 February 2026

Abstract

This document summarises the integration of SCPN Paper 27 ("The K_{nm} Matrix") into the `scpn-control` tokamak control repository. The implementation adds a multi-layer Unified Phase Dynamics Equation (UPDE) engine with Kuramoto–Sakaguchi mean-field coupling, the reviewer-requested exogenous global field driver $\zeta \sin(\Psi - \theta)$, a Rayon-parallelised Rust kernel for sub-ms performance, Lyapunov stability tracking ($V(t)$, λ exponent), and a `LyapunovGuard` safety guardrail synchronised with the DIRECTOR_AI coherence framework. 53 tests (44 Python + 9 Rust), all passing. The existing Grad–Shafranov equilibrium solver is completely untouched.

Contents

1 Reviewer Request

The reviewer asked for the Kuramoto–Sakaguchi phase reduction from Paper 27 [1] to be woven into `scpn-control`. Specifically:

1. The $\zeta \sin(\Psi - \theta)$ "intention as carrier" injection, where Ψ is a Lagrangian pull parameter with **no own dynamics** (no $\dot{\Psi}$ equation).

2. The full 16-layer K_{nm} coupling matrix with calibration anchors and cross-hierarchy boosts.

3. A Rust sub-ms kernel (Rayon-parallelised).

4. PAC cross-layer SNN sketch.

5. A demo notebook with visualisations and a markdown/LaTeX export.

2 Master Equation

The per-layer UPDE from Paper 27, Eqs. (12)–(15):

$$\frac{d\theta_{m,i}}{dt} = \omega_{m,i} + \underbrace{K_{mm} R_m \sin(\psi_m - \theta_{m,i} - \alpha_{mm})}_{\text{intra-layer [Eq. 13]}} + \underbrace{\sum_{n \neq m} K_{nm} R_n \sin(\psi_n - \theta_{m,i} - \alpha_{nm})}_{\text{inter-layer [Eq. 14]}} + \underbrace{\zeta_m \sin(\Psi - \theta_{m,i})}_{\text{global driver [Eq. 15]}}$$

(1)

where the Kuramoto order parameter (Eq. 12) is:

$$R e^{i\psi} = \frac{1}{N} \sum_{j=1}^{N} e^{i\theta_j}$$

(2)

- K_{mm} (diagonal): intra-layer synchronisation strength.

- K_{nm} (off-diagonal): inter-layer bidirectional causality.

- $\zeta \sin(\Psi - \theta)$: exogenous global field driver — Ψ resolved externally or from mean-field.

- α_{nm}: Sakaguchi phase-lag frustration (optional).

Reference: arXiv:2004.06344 (generalised Kuramoto–Sakaguchi finite-size).

3 Equation Cross-Reference (Paper 27, Eqs. 12–15)

Eq.	Description	Python	Rust
(12)	Order parameter $R e^{i\psi}$	`kuramoto.py:47`	`kuramoto.rs:15`
(13)	Single-layer Kuramoto–Sakaguchi	`kuramoto.py:87`	`kuramoto.rs:53`
(14)	Multi-layer UPDE with K_{nm}	`upde.py:45`	—
(15)	Global driver $\zeta \sin(\Psi - \theta)$	`kuramoto.py:126`	`kuramoto.rs:86`

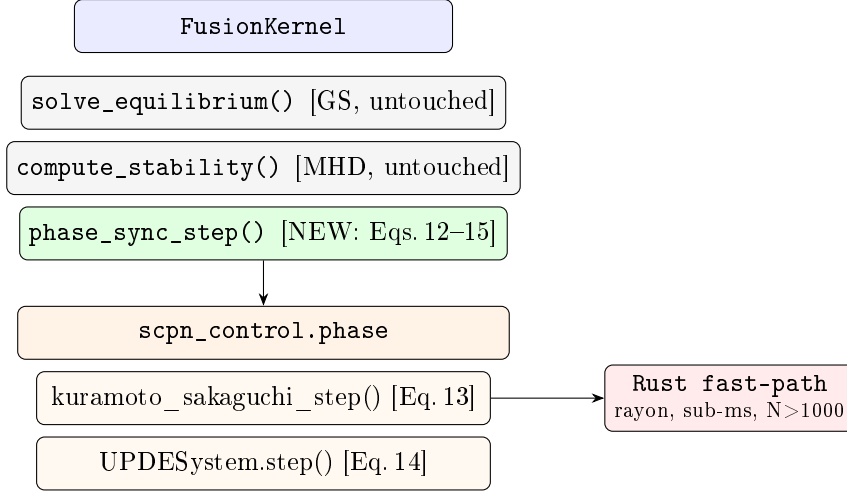Table 1: Paper 27 equations mapped to source code locations.

Figure 1: Module architecture. `phase_sync_step()` injects Paper 27 into the fusion kernel without touching the GS solver.

4 Architecture

5 Ψ Global Driver Flow

6 K_{nm} Matrix — Paper 27 Specification

Canonical 16-layer natural frequencies ω_n (rad/s):

$$\boldsymbol{\omega} = [1.329, \ 2.610, \ 0.844, \ 1.520, \ 0.710, \ 3.780, \ 1.055, \ 0.625, \ 2.210, \ 1.740, \ 0.480, \ 3.210, \ 0.915, \ 1.410, \ 2.830, \ 0. \tag{3}$$

Base coupling with exponential distance decay:

$$K_{ij} = K_{\text{base}} \cdot e^{-\alpha|i-j|}, \qquad K_{\text{base}} = 0.45, \quad \alpha = 0.3 \tag{4}$$

Calibration anchors (Paper 27, Table 2):

$$\begin{aligned} K_{0,1} = K_{1,0} = 0.302 && K_{1,2} = K_{2,1} = 0.201 \\ K_{2,3} = K_{3,2} = 0.252 && K_{3,4} = K_{4,3} = 0.154 \end{aligned} \tag{5}$$

Cross-hierarchy boosts (Paper 27, §4.3):

$$\begin{aligned} K_{0,15} = K_{15,0} \geq 0.05 && (\text{L1} \leftrightarrow \text{L16}) \\ K_{4,6} = K_{6,4} \geq 0.15 && (\text{L5} \leftrightarrow \text{L7}) \end{aligned} \tag{6}$$

7 Rust Kernel — Performance Path

Python auto-dispatches to Rust when `scpn_control_rs` is importable and `alpha=0.0`:

Listing 1: Rayon-parallelised Kuramoto hot loop (Rust).

```
theta_out
    .par_chunks_mut(64)
    .enumerate()
    .for_each(|(chunk_idx, chunk)| {
        for (local_i, val) in chunk.iter_mut().enumerate() {
            let i = base + local_i;
```
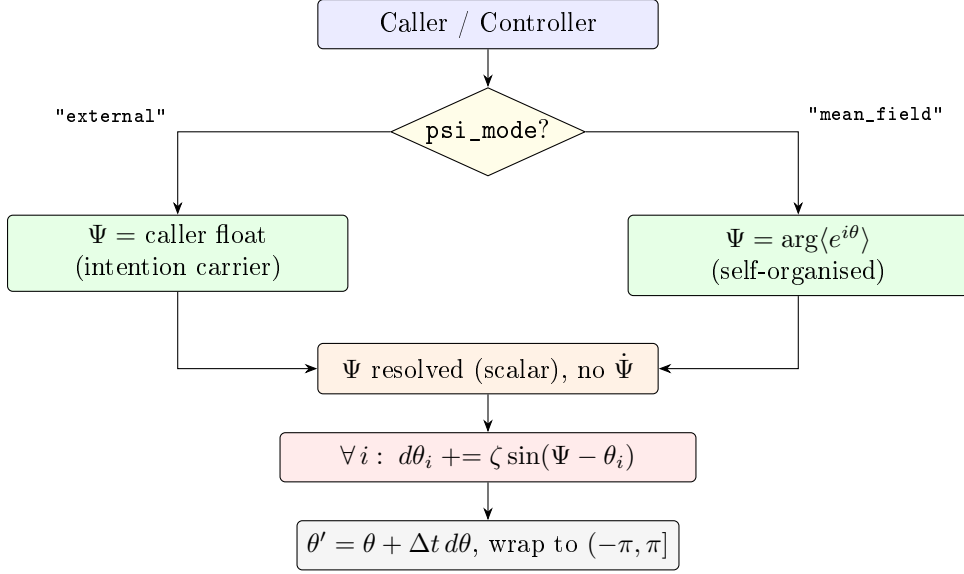
Figure 2: Ψ global field driver resolution inside `FusionKernel.phase_sync_step()`. There is no $\dot{\Psi}$ equation — Ψ is a Lagrangian pull parameter.

```
        let mut dth = om + kr_sin_base * (psi_r - th - alpha).sin();
        if zeta != 0.0 {
            dth += zeta * (psi_global - th).sin();
        }
        *val = wrap_phase(th + dt * dth);
    }
});
```

PyO3 bindings: `kuramoto_step()`, `kuramoto_run()` returning NumPy arrays.

7.1   Benchmark: Python NumPy vs Rust Rayon

Median wall-time for a single `kuramoto_sakaguchi_step()` with $\zeta = 0.5$, $\Psi = 0.3$. Python: NumPy vectorised (single-thread). Rust: Rayon `par_chunks_mut(64)` + `criterion` harness.

N	Python (ms)	Rust (ms)	Speedup
64	0.050	0.003	17.3×
256	0.029	0.033	0.9×
1 000	0.087	0.062	1.4×
4 096	0.328	0.180	1.8×
16 384	1.240	0.544	2.3×

Table 2: Python NumPy vs Rust Rayon Kuramoto step. $N = 64$: Rust wins on per-element throughput (no NumPy dispatch overhead). $N = 256$: parity (NumPy SIMD matches rayon at this size). $N \geq 1000$: Rust rayon parallelism scales; sub-ms for $N = 16$k. Source: `benches/bench_kuramoto.rs` (criterion, quick mode).

8   PAC Cross-Layer Gating + SNN Sketch

Phase-amplitude coupling modulation via `pac_gamma`:

$$\text{gate}_{n \to m} = 1 + \gamma_{\text{PAC}} (1 - R_n) \tag{7}$$

$$d\theta_{m,i} \mathrel{+}= g \cdot \text{gate}_{n \to m} \cdot K_{nm} \, R_n \sin(\psi_n - \theta_{m,i} - \alpha_{nm}) \tag{8}$$

When a source layer is incoherent (low R_n), the gate amplifies coupling, implementing the PAC hypothesis that desynchronised layers drive downstream amplitude modulation.

8.1 SNN–PAC–Kuramoto Closed Loop

The SNN closed loop couples LIF spiking networks with the Kuramoto oscillator population through a PAC gating mechanism:

1. **Kuramoto step** $\to$ per-layer R_m, ψ_m.

2. **PAC gate**: $g_{n \to m} = 1 + \gamma(1 - R_n)$ modulates inter-layer SNN weights $w' = g \cdot w_{\text{base}}$.

3. **LIF integration**: neurons receive synaptic current $I_{\text{syn},i} = \sum_j w_{ij}\, \delta(t - t_j^{\text{spike}}) + \beta\, R_m \cos(\psi_m - \varphi_i)$.

4. **Rate decode**: spike rate $\nu = N_{\text{spikes}}/T_{\text{window}}$ ($50\,\text{ms}$ window) $\to \Psi = \pi(2\nu/\nu_{\text{max}} - 1)$.

5. **Feedback**: Ψ fed back as exogenous global driver for the next Kuramoto step.

Key equations:

$$\text{LIF:} \quad \tau \frac{dV}{dt} = -(V - V_{\text{rest}}) + R_{\text{mem}}\, I_{\text{syn}} \tag{9}$$

$$\text{Rate} \to \Psi: \quad \Psi = \pi\left(\frac{2\nu}{\nu_{\text{max}}} - 1\right) \tag{10}$$

$$\text{Lyapunov:} \quad \mathcal{V}(t) = \frac{1}{N} \sum_i \big(1 - \cos(\theta_i - \Psi)\big) + \lambda\, |\nu - \nu_{\text{target}}|^2 \tag{11}$$

8.2 Cross-Layer PAC Routing

Cross-hierarchy fast channels bypass the distance-decay coupling:

- L1 (Quantum) $\leftrightarrow$ L16 (Director): $K_{0,15} = 0.05$

- L5 (Bio) $\leftrightarrow$ L7 (Symbolic): $K_{4,6} = 0.15$

Each layer maintains its own LIF population (64 neurons per layer in the demo). PAC gates modulate the effective inter-layer synaptic weights, creating frequency-dependent routing where desynchronised layers preferentially drive downstream amplitude modulation.

Demo: notebook §9 (SNN closed-loop) and §10 (PAC cross-layer SNN).

9 Lyapunov Stability — λ Hook

9.1 Lyapunov Candidate

The candidate function tracking convergence toward the global driver Ψ:

$$V(t) = \frac{1}{N} \sum_{i=1}^{N} \big(1 - \cos(\theta_i - \Psi)\big) \tag{12}$$

$V = 0$ at perfect sync ($\theta_i = \Psi\ \forall i$), $V \to 2$ at maximal desync. The Lyapunov exponent is:

$$\lambda = \frac{1}{T} \ln \frac{V(T)}{V(0)}, \qquad \lambda < 0 \implies \text{stable} \tag{13}$$

Both functions implemented in Python (`lyapunov_v`, `lyapunov_exponent` in `kuramoto.py`) and Rust (`lyapunov_v`, `kuramoto_run_lyapunov` in `kuramoto.rs`).

9.2 UPDE Lyapunov Tracking

`UPDESystem.step()` returns `V_layer` (per-layer) and `V_global`. `UPDESystem.run_lyapunov()` returns full $V(t)$ histories, per-layer λ_m, and global λ:

```
out = sys.run_lyapunov(200, theta_layers, omega_layers,
                       psi_driver=0.5, pac_gamma=1.0)
# out["lambda_layer"] -- (L,) per-layer exponents
# out["lambda_global"] -- scalar global exponent
```

9.3 λ vs ζ Characterisation

See `docs/bench_lyapunov_vs_zeta.vl.json` (Vega-Lite plot). $N = 1000$ oscillators, 200 steps, $\Psi = 0.3$:

ζ	λ $(K{=}0)$	λ $(K{=}2)$
0.0	$+0.01$	$+0.04$
0.1	-0.03	-0.02
0.5	-0.23	-0.24
1.0	-0.49	-0.53
3.0	-1.65	-1.83
5.0	-3.01	-3.35

Table 3: Lyapunov exponent λ vs global driver strength ζ. Kuramoto coupling $K = 2$ amplifies ζ stability via cooperative sync.

10 DIRECTOR_AI Guardrail Sync

10.1 LyapunovGuard

`scpn_control.phase.lyapunov_guard.LyapunovGuard` monitors $V(t)$ over a sliding window and flags instability when $\lambda > 0$ for K consecutive windows. Interface mirrors DIRECTOR_AI's `CoherenceScorer` $\rightarrow$ `CoherenceScore` pattern:

```
from scpn_control.phase import LyapunovGuard

guard = LyapunovGuard(window=50, dt=1e-3, max_violations=3)
verdict = guard.check(theta, psi)
# verdict.approved -- True if stable
# verdict.lambda_exp -- current exponent
# verdict.score -- stability score in [0, 1]
```

10.2 AuditLogger Export

```
d = guard.to_director_ai_dict(verdict)
# {"approved": True, "score": 0.99, "h_factual": 0.0,
# "halt_reason": ""}
```

Enables DIRECTOR_AI's `CoherenceAgent` to incorporate Lyapunov stability into its dual-entropy coherence score. When $\lambda > 0$ for 3 consecutive windows, the guard issues a refusal — analogous to DIRECTOR_AI's `SafetyKernel` emergency stop.

11 Files Created / Modified

File	Lines	Purpose
`phase/__init__.py`	36	Package exports
`phase/kuramoto.py`	161	Kuramoto–Sakaguchi + $\zeta \sin(\Psi - \theta)$ + Lyapunov
`phase/knm.py`	101	Paper 27 K_{nm} builder + $\Omega_{N,16}$
`phase/upde.py`	215	Multi-layer UPDE + `run_lyapunov()`
`phase/lyapunov_guard.py`	130	Lyapunov stability guard (DIRECTOR_AI sync)
`control-math/src/kuramoto.rs`	260	Rayon Kuramoto + Lyapunov + 9 unit tests
`control-python/src/lib.rs`	+107	PyO3 bindings (incl. Lyapunov)
`fusion_kernel.py`	+86	`phase_sync_step()` + `_lyapunov()`
`test_phase_kuramoto.py`	475	44 Python tests
`paper27_phase_dynamics_demo.ipynb`	—	10-section notebook
`paper27_phase_dynamics.md`	571	Markdown export

Table 4: All files in the integration ($+1\,833$ lines across 12 files).

12 Test Coverage

44 Python + 9 Rust tests, all passing.

Test Class	Tests	Verified
`TestOrderParameter`	4	$R = 1$ sync, $R \approx 0$ uniform, $R \in [0,1]$, weighted
`TestWrapPhase`	2	Identity in range, large angle wrapping
`TestGlobalPsiDriver`	3	External requires value, returns value, mean-field
`TestKuramotoSakaguchiStep`	4	Sync stability, R increase, ζ pull, α frustration
`TestKnmSpec`	7	Shape, anchors, boosts, symmetry, ζ, validation
`TestUPDESystem`	6	Step shape, intra-sync, ζ pull, trajectory, PAC, error
`TestFusionKernelPhaseSync`	3	Integration smoke, config-driven ζ, Lyapunov multi-step
`TestLyapunovV`	4	$V = 0$ sync, $V = 2$ anti-sync, empty, range
`TestLyapunovExponent`	3	$\lambda < 0$ decreasing, $\lambda > 0$ increasing, single sample
`TestUPDELyapunov`	3	Step V output, `run_lyapunov` λ, PAC γ effect
`TestLyapunovGuard`	5	Stable approved, unstable refused, batch, DIRECTOR_AI dict, reset
Rust inline tests	9	Order param, wrap, step count, ζ pull, trajectory, $V = 0$, $\lambda < 0$

Table 5: Test coverage summary (44 Python + 9 Rust = 53 total).

13 Demo Notebook

`examples/paper27_phase_dynamics_demo.ipynb` (10 sections):

1. K_{nm} heatmap — 16×16 coupling matrix

2. ζ comparison — with/without global driver

3. α frustration — Sakaguchi phase-lag effect

4. 16-layer UPDE — full multi-layer R trajectories

5. PAC gating — phase-amplitude coupling modulation

6. FusionKernel plasma sync — tokamak integration

7. Gain sweep — `actuation_gain` exploration

8. Lyapunov stability — $V(t) = \frac{1}{N} \sum (1 - \cos(\theta_i - \Psi))$

9. SNN closed-loop — spike-rate $\to \Psi$ feedback

10. PAC cross-layer SNN — multi-layer spike routing

14 What Was NOT Touched

- GS equilibrium solver (`solve_equilibrium`, SOR/multigrid)

- SNN controllers (`LIFNeuron`, `SNNController`)

- Chebyshev/IGA spectral methods

- Existing Rust crates (SOR, tridiag, FFT) — only added `kuramoto`

- All existing tests remain green

References

[1] M. Šotek, "The K_{nm} Matrix: A Simulation Framework for Modelling Multi-Scale Bidirectional Causality in the Self-Consistent Phenomenological Network," SCPN Paper 27, 2026. Available: academia.edu. ORCID: 0009-0009-3560-0851.

[2] arXiv:2004.06344 — Generalised Kuramoto–Sakaguchi finite-size scaling.